

# File Systems

2024 Semester 2 COMPSCI 340: Operating Systems  
Talía Xu

Lecture 1  
1.0.0

## **I'm Talia, your instructor**

I am a lecturer in the Computer Systems Group.

I am about to receive my PhD from TU Delft.

I previously went to University of Toronto for BASc and MASc.

## **I'm Talia, your instructor**

I worked at a few industry places, mostly as an intern

- Canada: IBM, Arista Networks, Amazon
- US: Reach Power (startup)
- UK: Nokia Bell Labs

My research interest is

- Using light for communication and sensing
- Novel sensing systems for health / mental well-beings
- Generally involves some hardware or embedded systems

## **I'm Talia, your instructor**

Office: 303s-592

Office hours: Mondays & Fridays 1 pm - 2 pm

Email: [talia.xu@auckland.ac.nz](mailto:talia.xu@auckland.ac.nz)

## **Please provide feedback**

Let me know what you like, dislike, or want to see more of

I'm open to suggestions!

## **These books complement lectures**

“Operating Systems: Three Easy Pieces”

by Remzi Arpaci-Dusseau and Andrea Arpaci-Dusseau

“Modern Operating Systems”

by Andrew S. Tanenbaum

Test: You are only responsible for what have been covered in class.

## Why should we care about OS?

Distinguished Lecture by Amin Vahdat

- **End of Moore's Law:** The traditional approach of relying on exponential hardware improvement is no longer sustainable.
- **Shifting Demands:** The growing demand for machine learning and data processing requires a different approach to computing infrastructure.
- **System-Level Optimization:** Optimizing the entire system, rather than individual components, can unlock greater gains in capacity and capability.

## **File System: How to manage a persistent device?**

What are the APIs?

What are the important aspects of the implementation?



## File system: How to manage a persistent device?

You have a **1 MB file** named foo stored in the directory /bar

- What is the **full path name** of the file?

How are files organized on the disk?


- What **data structures** are used to store the data of the file?
- The **metadata**?
- The **location**?

How do you access data on the disk?

- When you open/write/read the file, what **function calls** are evoked?
- If we write another 1MB data to /bar/foo, where does this data go?
- Which structures on the disks are read from and written to?

## File system is an abstract concept

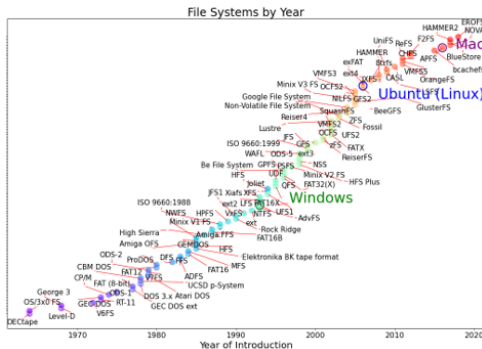
A File system describes a way of organizing and storing files on a storage device.



What are some requirements of a file system?

- **Availability:** Ensure data can be accessed and used reliably.
- **Permanence:** Store data permanently (or an approximation of it).
- **Concurrent Access:** Make data sharable with other programs or users

## File system implementations are design choices



Some popular file system implementations:

- Unix (Linux and Mac)
- Windows
- Specialized system (Mainframe)

## **What is a file?**


Textbook: A named collection of related information that is recorded on secondary storage.

For most users, the basic unit of interaction with a file system is a file.

- Is a directory a file?
- Is a process a file?
- Is a device a file?

## Usual layout of a POSIX Filesystem

POSIX is a set of standards that provides a common interface for operating systems.



Working Directory: /home/talia

- What is the absolute and relative path to **todo.txt**? To **usb**?

## POSIX Filesystem

todo.txt Relative: ./todo.txt

todo.txt Absolute: /home/jon/todo.txt

usb Relative: ../../mnt/usb

usb Absolute: /mnt/usb

Special symbols:

- . — Current directory

- .. — Parent directory

- ~ — User's home directory (\$HOME)

Relative paths are calculated from current working directory (\$PWD)

## **You Can Access Files Sequentially or Randomly**

### **Sequential access**

- Each read advances the position inside the file
- Writes are appended and the position set to the end afterwards

### **Random access**

- Records can be read/written to the file in any order
- A specific position is required for each operation

## POSIX Filesystem

```
int open(const char *pathname, int flags, mode_t mode);

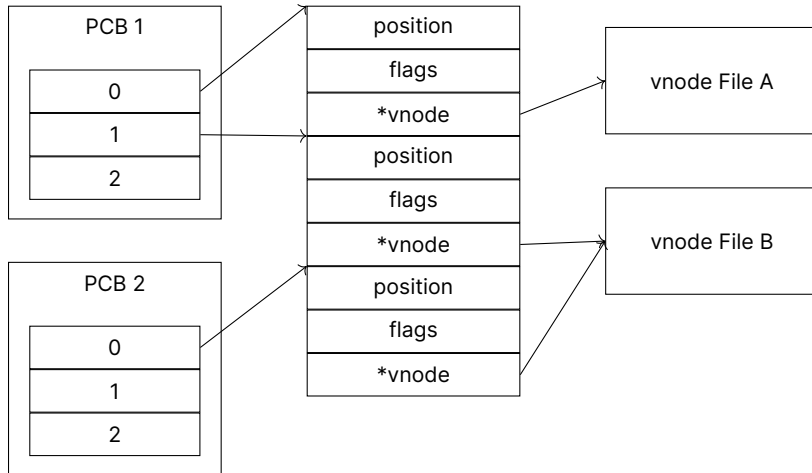
// flags can specify which operations: O_RDONLY, O_WRONLY, O_RDWR
// also: O_APPEND moves the position to the end of the file initially

off_t lseek(int fd, off_t offset, int whence);

// lseek changes the position to the offset
// whence can be one of: SEEK_SET, SEEK_CUR, SEEK_END
//   set makes the offset absolute, cur and end are both relative
```



## File Tables Are Stored in the Process Control Block (PCB)



## Each Process Contains a File Table in its PCB

A **Fie Descriptor** is an **index** in the table

Each item points to a system-wide *global open file* (**GOF**) table

The GOF table holds information about the **seek position** and **flags**

- It also points to a *VNode* (supports read/write/etc)

A **vnode** (virtual mode) holds information about the file


- vnodes can represent regular files, pipes, network sockets, etc.

## **What happens in a fork?**

PCB is copied – local open file table gets inherited

Both PCBs point to the same Global Open Table entry

## Both Processes Point to the Same GOF Entry

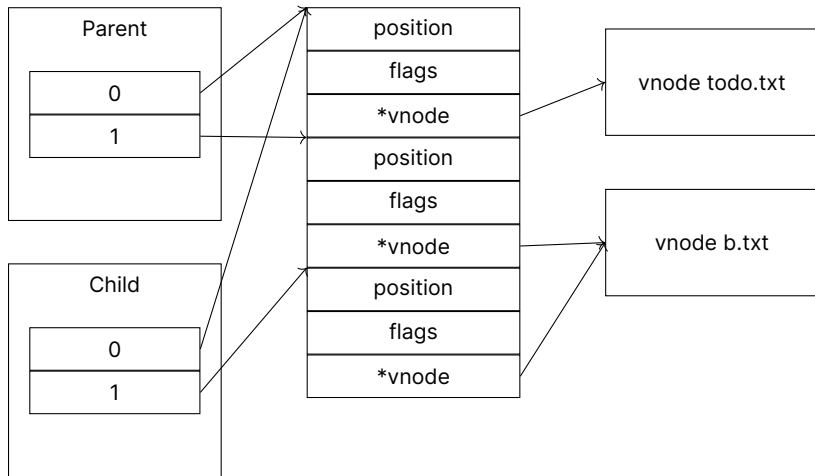


## How many LOF and GOF Entries Exist? What is the Relationship?

```
open("todo.txt", O_RDONLY);  
fork();  
open("b.txt", O_RDONLY);
```

Assume there are no previously opened files (not even the standard ones)

## There are 2 LOF Entries Each, and 3 GOF Entries



## What happens in Fork?

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main() {
    FILE *file;
    file = fopen("test.txt", "w");

    // Fork the process
    pid_t pid = fork();

    fputs("hello\n", file);
    fputs("hello again\n", file);

    return 0;
}
```

## What happens in Fork?

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main() {
    FILE *file;
    file = fopen("test.txt", "w");

    fputs("hello\n", file);
    fputs("hello again\n", file);

    // Fork the process
    pid_t pid = fork();

    return 0;
}
```



## What happens in Fork?

```
int main() {  
    FILE *file;  
    file = fopen("test.txt", "r");  
  
    pid_t pid = fork();  
  
    if (pid == 0) {  
        char line[256];  
        if (fgets(line, sizeof(line), file) != NULL) {  
            printf("Child read: %s", line);  
        }  
    } else {  
        char line[256];  
        if (fgets(line, sizeof(line), file) != NULL) {  
            printf("Parent read: %s", line);  
        }  
    }  
}
```

## Reading

Textbook

- 13.1 File Concept
- 13.2 Access Methods