

Advanced Scheduling

2025 Semester 1 SOFTENG 370: Operating Systems

Talia Xu

Lecture 8

2.0.0

How to schedule without perfect knowledge?

How can we design a scheduler system that both

- minimizes response time for interactive jobs
- minimizes turnaround time without a priori knowledge of job length


We may favor some processes over others.

Multi-level Feedback Queue (MLFQ)

There are multiple queues with different priority levels.

- If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).
- If $\text{Priority}(A) = \text{Priority}(B)$, A and B run in RR.

A Snapshot of MLFQ



Dynamic Priority

We may lead processes to *starvation* if there's a lot of higher priority processes

One solution is to have the OS dynamically change the priority

Dynamic Priority Attempt 1


New jobs are placed at the top queue with highest priority.

Allotment: Each job is allocated a time slice before the scheduler reduces its priority.


If a job uses up its allotment while running, its priority is reduced.

A job's usage resets if it gives up the CPU before its allotted time is used up.

Dynamic Priority Attempt 1



Dynamic Priority Attempt 1



Dynamic Priority Attempt 1

Prioritizes I/O-intensive interactive jobs (good response time).


Can cause starvation.

Users can try to cheat the system.

What if a long-running program becomes interactive later?

Dynamic Priority Attempt 2

After some time period S , move all the jobs in the system to the topmost queue.



Dynamic Priority Attempt 3

A job's usage **does not reset** if it gives up the CPU before its allotted time is used up.

How many queues?

How long is an allotment?

How long is the time period S ?

Scheduling Can Get Complicated

There's no "right answer", only trade-offs

We haven't talked about multiprocessor scheduling yet

We'll assume symmetric multiprocessing (SMP)

- All CPUs are connected to the same physical memory

- The CPUs have their own private cache (at least the lowest levels)

One Approach is to Use the Same Scheduling for All CPUs

There's still only one scheduler

It just keeps adding processes while there's available CPUs

Advantages

Good CPU utilization

Fair to all processes

Disadvantages

Not scalable (everything blocks on global scheduler)

Poor cache locality

This was the approach in Linux 2.4

We Can Create Per-CPU Schedulers

When there's a new process, assign it to a CPU

One strategy is to assign it to the CPU with the lowest number of processes

Advantages

- Easy to implement

- Scalable (there's no blocking on a resource)

- Good cache locality

Disadvantages

- Load imbalance

- Some CPUs may have less processes, or less intensive ones

We Can Compromise between Global and Per-CPU

- Keep a global scheduler that can rebalance per-CPU queues

 - If a CPU is idle, take a process from another CPU (work stealing)

- You may want more control over which processes can switch

 - Some may be more sensitive to caches

- Use *processor affinity*

 - The preference of a process to be scheduled on the same core

- This is a simplified version of the O(1) scheduler in Linux 2.6

Another Strategy is “Gang” Scheduling

Multiple processes may need to be scheduled simultaneously

The scheduler on each CPU cannot be completely independent

“Gang Scheduling” (Coscheduling)

Allows you to run a set of processes simultaneously (acting as a unit)

This requires a global context-switch across all CPUs

Real-Time Scheduling is Yet Another Problem

Real-time means there are time constraints, either for a deadline or rate
e.g. audio, autopilot

A hard real-time system

Required to guarantee a task completes within a certain amount of time

A soft real-time system

Critical processes have a higher priority and the deadline is met in practice

Linux is an example of soft real-time

Linux Also Implements FCFS and RR Scheduling

You can search the source tree: FCFS (SCHED_FIFO) and RR (SCHED_RR)

Use a multilevel queue scheduler for processes with the same priority

Also let the OS dynamically adjust the priority

Soft real-time processes:

Always schedule the highest priority processes first

Normal processes:

Adjust the priority based on aging

Real-Time Processes Are Always Prioritized

The soft real-time scheduling policy will either be `SCHED_FIFO` or `SCHED_RR`

There are 100 static priority levels (0—99)

Normal scheduling policies apply to the other processes (`SCHED_NORMAL`)


By default the priority is 0

Priority ranges from $[-20, 19]$

Processes can change their own priorities with system calls:

`nice`, `sched_setscheduler`

Linux Priority Unifies Soft Real-time and Normal Processes



Linux Scheduler Evolution

2.4—2.6, a $O(N)$ global queue

Simple, but poor performance with multiprocessors and many processes

2.6—2.6.22, a per-CPU run queue, $O(1)$ scheduler

Complex to get right, interactivity had issues

No guarantee of fairness

2.6.23—Present, the completely fair scheduler (CFS)

Fair, and allows for good interactivity

The $O(1)$ Scheduler Has Issues with Modern Processes

Foreground and background processes are a good division

Easier with a terminal, less so with GUI processes

Now the kernel has to detect interactive processes with heuristics

Processes that sleep a lot may be more interactive

This is ad hoc, and could be unfair

How would we introduce fairness for different priority processes?

Use different size time slices

The higher the priority, the larger the time slice

There are also situations where this ad hoc solution could be unfair


Ideal Fair Scheduling

Assume you have an infinitely small time slice
If you have n processes, each runs at $\frac{1}{n}$ rate

1 Process



3 Processes



CPU usage is divided equally among every process

Example IFS Scheduling

Consider the following processes:

Process	Arrival Time	Burst Time
P ₁	0	8
P ₂	0	4
P ₃	0	16
P ₄	0	4

Assume that each vertical slice can execute 4 time units.
Each box represents the time units spend executing

	0				16		24		32
P ₁	1	2	3	4	6	8			
P ₂	1	2	3	4					
P ₃	1	2	3	4	6	8	12	16	
P ₄	1	2	3	4					

IFS is the Fairest but Impractical Policy

This policy is fair, every process gets an equal amount of CPU time
Boosts interactivity, has the ideal response time

However, this would perform way too many context switches

You have to constantly scan all processes, which is $O(N)$

Completely Fair Scheduler (CFS)

For each runnable process, assign it a “virtual runtime”

At each scheduling point where the process runs for time t

Increase the virtual runtime by $t \times \text{weight}$ (based on priority)

The virtual runtime monotonically increases

Scheduler selects the process based on the lowest virtual runtime

Compute its dynamic time slice based on the IFS

Allow the process to run, when the time slice ends repeat the process

CFS is Implemented with Red-Black Trees

- A red-black tree is a self-balancing binary search tree

 - Keyed by virtual runtime

 - $O(\lg N)$ insert, delete, update, find minimum

- The implementation uses a red-black tree with nanosecond granularity

 - Doesn't need to guess the interactivity of a process

- CFS tends to favour I/O bound processes by default

 - Small CPU bursts translate to a low virtual runtime

 - It will get a larger time slice, in order to catch up to the ideal